

Operating System Project 3

1. Revise Source Codes

```


1507 int filemap_fault(struct vm_area_struct *vma, struct vm_fault *vmf)
1508 {
1509     int error;
1510     struct file *file = vma->vm_file;
1511     struct address_space *mapping = file->f_mapping;
1512     struct file_ra_state *ra = &file->f_ra;
1513     struct inode *inode = mapping->host;
1514     pgoff_t offset = vmf->pgoff;
1515     struct page *page;
1516     pgoff_t size;
1517     int ret = 0;
1518
1519     /* OS P33: Instrument message */
1520     printk(KERN_CRIT "%s, %X\n", current->comm, vmf->virtual_address);
1521
1522     size = (i_size_read(inode) + PAGE_CACHE_SIZE - 1) >> PAGE_CACHE_SHIFT;
1523     if (offset >= size)
1524         return VM_FAULT_SIGBUS;
1525
1526 +--- 3 lines: Do we have something in the page cache already?-----
1529     page = find_get_page(mapping, offset);
1530     if (likely(page)) {
1531         /*
1532          * We found the page, so try async readahead before
1533          * waiting for the lock.
1534          */
1535         /* OS P33: cancel readahead policy */
1536         do_async_mmap_readahead(vma, ra, file, page, offset);
1537         lock_page(page);
1538
1539         /* Did it get truncated? */
1540         if (unlikely(page->mapping != mapping)) {
1541             unlock_page(page);
1542             put_page(page);
1543             goto no_cached_page;
1544         }
1545     } else {
1546         /* No page in the page cache at all */
1547         /* OS P33: cancel readahead policy */
1548         do_sync_mmap_readahead(vma, ra, file, offset);
1549         count_vm_event(PGMAJFAULT);
1550         ret = VM_FAULT_MAJOR;

```

2. Trace Codes

2.1. filemap_fault()

The function will be called when page fault occurs. First, it checks whether the target page is in the page cache, if not, then it use goto method to jump to no_cached_page code session to read pages from storage. if the page is in the page cache, it checks if the page is truncated, if not, it will samely jump to no_cached_page session. After going through the control flow mentioned before, it then consequently goes to retry_find session. The last step is checking whether the page is up-to-date, if yes, the function directly returns, otherwise, it jumps to page_no_uptodate session to update the page before returning.



2.2. mmap()

The mmap() function shall establish a mapping between a process' address space and a file, shared memory object, or typed memory object.

3. Compare Pure Demand Paging & Read Ahead Algorithm

3.1. Test Outputs

pure demand paging : (pd_output)

```
7168 102010677
7169 671903510
7170 1608468492
7171 # of major pagefault: 6567
7172 # of minor pagefault: 239
7173 # of resident set size: 26676 KB
```

read ahead algorithm : (ra_output)

```
7168 102010677
7169 671903510
7170 1608468492
7171 # of major pagefault: 4201
7172 # of minor pagefault: 2605
7173 # of resident set size: 26680 KB
```

3.2. Comparison Table

	major pagefaults	minor pagefaults	resident set size
pure demand paging	6567	239	26676 KB
read ahead algoritm	4201	2605	26680 KB